

## Problem 8: UK Election Results, 2015-2019

Rakesh Vishwabrahmana

```
# Require Packages
library(tidyverse)

## -- Attaching packages ----- tidyverse
1.3.1 --

## v ggplot2 3.3.5      v purrr  0.3.4
## v tibble  3.1.5      v dplyr  1.0.7
## v tidyr   1.1.4      v stringr 1.4.0
## v readr   2.0.2      v forcats 0.5.1

## -- Conflicts -----
tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(GGally)

## Warning: package 'GGally' was built under R version 4.1.2

## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2

library(ggfortify)
library(plyr)

## -----
## -----

## You have loaded plyr after dplyr - this is likely to cause problems.
## If you need functions from both plyr and dplyr, please load plyr first,
## then dplyr:
## library(plyr); library(dplyr)

## -----
## -----

##
## Attaching package: 'plyr'

## The following objects are masked from 'package:dplyr':
##
##   arrange, count, desc, failwith, id, mutate, rename, summarise,
##   summarize
```

```
## The following object is masked from 'package:purrr':
##
## compact
library(olsrr)
## Warning: package 'olsrr' was built under R version 4.1.2
##
## Attaching package: 'olsrr'
## The following object is masked from 'package:datasets':
##
## rivers
library(car)
## Warning: package 'car' was built under R version 4.1.2
## Loading required package: carData
##
## Attaching package: 'car'
## The following object is masked from 'package:dplyr':
##
## recode
## The following object is masked from 'package:purrr':
##
## some
```

## Question: 1

### Solution:

#### (a)

In the following r chunk, we load all relevant CSV files into R data frame variables and convert string variable as factor.

```
GE2015 <- read_csv("GE2015-results.csv")
## Rows: 650 Columns: 28
## -- Column specification -----
## Delimiter: ","
## chr (11): ons_id, ons_region_id, constituency_name, county_name,
region_name...
## dbl (17): electorate, valid_votes, invalid_votes, majority, con, lab, ld,
uk...
```

```

##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this
message.

GE2015 <- GE2015%>%
  mutate_if(is.character,as.factor)

GE2017 <- read_csv("GE2017-results.csv")

## Rows: 650 Columns: 29

## -- Column specification -----
## Delimiter: ","
## chr (10): ons_id, ons_region_id, constituency_name, county_name,
region_nam...
## dbl (18): electorate, valid_votes, invalid_votes, majority, con, lab, ld,
u...
## dtm (1): declaration_time

##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this
message.

GE2017 <- GE2017%>%
  mutate_if(is.character,as.factor)

GE2019 <- read_csv("GE2019-results.csv")

## Rows: 650 Columns: 32

## -- Column specification -----
## Delimiter: ","
## chr (13): ons_id, ons_region_id, constituency_name, county_name,
region_nam...
## dbl (18): electorate, valid_votes, invalid_votes, majority, con, lab, ld,
b...
## dtm (1): declaration_time

##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this
message.

GE2019 <- GE2019 %>%
  mutate_if(is.character,as.factor)

demographics <- read_csv("demographics.csv")

```

```
## Rows: 650 Columns: 13

## -- Column specification -----
-----
## Delimiter: ","
## chr (2): ons_id, constituency
## dbl (11): income, age.0.15, age.65.over, foreignborn, employment,
outofwork,...

##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this
message.

demographics <- demographics %>%
  mutate_if(is.character, as.factor)
```

### (b)

We add an extra categorical column **year** using the year of the election 2015, 2017 and 2019 to each results data frame.

```
GE2015 <- GE2015 %>%
  mutate(year= "2015")
GE2017 <- GE2017 %>%
  mutate(year= "2017")
GE2019 <- GE2019 %>%
  mutate(year= "2019")
```

### (c)

We will merge all the three separate year-specific results data frames into a single data frame containing all results in the following chunk.

```
mrg_data <- plyr::rbind.fill(GE2015, GE2017, GE2019)
dim(mrg_data)

## [1] 1950 34
```

### (d)

In the given chunk below we merge the combined results data frame together with the demographics data to produce a single data frame.

```
full_data <- left_join(mrg_data, demographics, by = "ons_id", copy=TRUE)
```

### (e)

We calculate the conservative vote share, then apply a logit transform to this value and create a new column named **con\_share** of the data frame.

```
full_data <- full_data %>%
  mutate(con_share = car::logit(con/valid_votes))
```

```
## Warning in car::logit(con/valid_votes): proportions remapped to (0.025,
0.975)
```

(f)

We select only relevant columns for the modeling purpose which will contain response and explanatory variable and define the object as final data frame.

```
new_data <- full_data %>%
  select(c(year, country_name, income, age.0.15, age.65.over,
    foreignborn, employment, outofwork, white, commute.car,
    commute.bike, health.good, health.bad, con_share))
```

(g)

We will filter out rows associated with *Northern Ireland* constituencies from the data set.

```
mod_data <- new_data %>%
  filter(country_name != "Northern Ireland") %>%
  mutate_if(is.character, as.factor) %>%
  drop_na()
```

```
dim(mod_data)
```

```
## [1] 1896 14
```

(h)

In the following chunk, we will use base function `summary()` to summarize variables in our final data frame.

```
summary(mod_data)
```

```
##      year      country_name      income      age.0.15
## 2015:632  England      :1599  Min.    :17900  Min.    :11.30
## 2017:632  Northern Ireland:  0  1st Qu.:21200  1st Qu.:17.40
## 2019:632  Scotland      : 177  Median :22600  Median :18.40
##              Wales      : 120  Mean     :23423  Mean     :18.63
##              3rd Qu.:24900  3rd Qu.:19.80
##              Max.      :44300  Max.      :30.40
##      age.65.over  foreignborn  employment  outofwork
## Min.    : 5.5    Min.      : 2.00  Min.      :42.00  Min.      : 0.480
## 1st Qu.:14.6    1st Qu.: 4.60  1st Qu.:58.88  1st Qu.: 1.417
## Median :16.8    Median : 7.50  Median :62.20  Median : 2.320
## Mean     :16.8    Mean     :11.85  Mean      :61.78  Mean      : 2.627
## 3rd Qu.:19.4    3rd Qu.:13.57  3rd Qu.:65.80  3rd Qu.: 3.475
## Max.      :32.2    Max.      :59.30  Max.      :74.60  Max.      :10.250
##      white      commute.car  commute.bike  health.good
## Min.      :23.10  Min.      :10.61  Min.      : 0.200  Min.      :70.60
## 1st Qu.:85.95    1st Qu.:62.41  1st Qu.: 1.238  1st Qu.:78.70
## Median :94.40    Median :69.44  Median : 1.960  Median :81.35
## Mean      :88.04  Mean      :65.03  Mean      : 2.602  Mean      :81.11
```

```
## 3rd Qu.:97.60 3rd Qu.:74.46 3rd Qu.: 3.125 3rd Qu.:83.50
## Max. :99.20 Max. :83.33 Max. :30.350 Max. :89.00
## health.bad con_share
## Min. : 2.400 Min. :-3.6636
## 1st Qu.: 4.475 1st Qu.: -0.9117
## Median : 5.300 Median :-0.2304
## Mean : 5.667 Mean :-0.4054
## 3rd Qu.: 6.600 3rd Qu.: 0.1800
## Max. :11.600 Max. : 1.1194
```

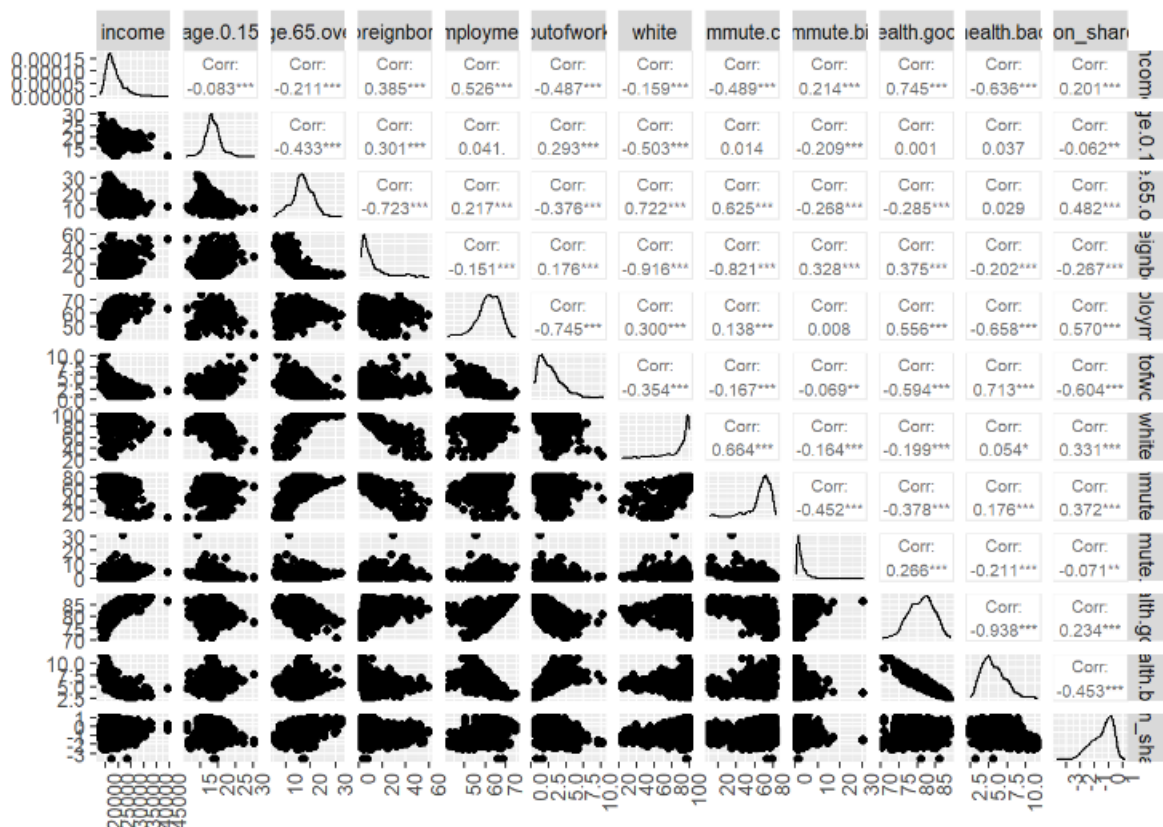
## Question: 2

### Solution:

The following R chunk below we will create a data frame with all the numeric variables and visualize there relationship using a pair matrix plot.

```
num_data <- mod_data %>%
  select_if(is.numeric)

ggpairs(num_data, upper = list(continuous = wrap("cor", size = 2.5))) +
  theme(axis.text.x = element_text(angle = 90, hjust = 1))
```



### Question: 3

#### Solution:

In the given chunk below we will give a table which will show whether the relationship with the response variable is positively correlated (+), not significantly correlated (o), or negatively correlated (-).

```
pos <- c("income", "age.65.over", "employment", "white", "commute.car",
"health.good")
nsig <- rep(NA, 6)
neg <- c("age.0.15", "foreignborn", "outofwork", "commute.bike", "health.bad",
NA)
cor_tab <- data.frame(pos, nsig, neg)
names(cor_tab)<- c("+", "o", "-")
cor_tab

##           +  o      -
## 1      income NA   age.0.15
## 2 age.65.over NA  foreignborn
## 3  employment NA   outofwork
## 4         white NA commute.bike
## 5 commute.car NA   health.bad
## 6 health.good NA          <NA>
```

### Question: 4

#### Solution:

In the following chunk below we will find the association between the response and categorical variables and find the group that predicts the highest and lowest value of the response variable.

```
cat_data <- mod_data %>%
  select(con_share, year, country_name)
mod1 <- aov(con_share~ year+country_name, data = cat_data)
summary(mod1)

##           Df Sum Sq Mean Sq F value Pr(>F)
## year           2    39.2   19.59   45.23 <2e-16 ***
## country_name    2   179.6   89.81  207.37 <2e-16 ***
## Residuals  1891   819.0    0.43
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

cat_data%>%
  group_by(year) %>%
  summarise(ave_con_share=mean(con_share))

##   ave_con_share
## 1    -0.4054456
```

```

cat_data%>%
  group_by(country_name) %>%
  summarise(ave_con_share=mean(con_share))

##   ave_con_share
## 1      -0.4054456

data.frame(Highest = c("2019", "England"), lowest=c("2015", "Scotland"))

##   Highest  lowest
## 1    2019    2015
## 2 England Scotland

```

## Question: 5

### Solution:

```

# (a)
null_mod <- lm(con_share~1,data=mod_data)
# (b)
full_mod <- lm(con_share~.,data=mod_data)
# (c)
int.year_mod <- lm(con_share~.+year*.,data=mod_data)
# (d)
int.all_mod <- lm(con_share~*.,data=mod_data)
# (e)
step_c <- step(int.year_mod, trace=0)
# (f)
step_d <- step(int.all_mod, trace=0)

```

## Question: 6

### Solution:

#### (a) Degrees of error freedom

```

df_e <- c(null_mod$df.residual,
          full_mod$df.residual,
          int.year_mod$df.residual,
          int.all_mod$df.residual,
          step_c$df.residual,
          step_d$df.residual)

```

#### (b) Degrees of model freedom

```

N <- dim(mod_data)[1]
df_mod <- c(N - null_mod$df.residual,
            N - full_mod$df.residual,
            N - int.year_mod$df.residual,
            N - int.all_mod$df.residual,
            N - step_c$df.residual,
            N - step_d$df.residual)

```



### (c) Multiple R-squared

```
r2 <- c(summary(null_mod)$r.squared,
        summary(full_mod)$r.squared,
        summary(int.year_mod)$r.squared,
        summary(int.all_mod)$r.squared,
        summary(step_c)$r.squared,
        summary(step_d)$r.squared)
```

### (d) Adjusted R-squared

```
adj_r2 <- c(summary(null_mod)$adj.r.squared,
            summary(full_mod)$adj.r.squared,
            summary(int.year_mod)$adj.r.squared,
            summary(int.all_mod)$adj.r.squared,
            summary(step_c)$adj.r.squared,
            summary(step_d)$adj.r.squared)
```

### (e) Residual standard error

```
se_e <- c(summary(null_mod)$sigma,
          summary(full_mod)$sigma,
          summary(int.year_mod)$sigma,
          summary(int.all_mod)$sigma,
          summary(step_c)$sigma,
          summary(step_d)$sigma)
```

### (f) Akaike information criterion

```
aic <- c(AIC(null_mod),
        AIC(full_mod),
        AIC(int.year_mod),
        AIC(int.all_mod),
        AIC(step_c),
        AIC(step_d))
```

### (h) PRESS

```
press <- c(olsrr::ols_press(null_mod),
          olsrr::ols_press(full_mod),
          olsrr::ols_press(int.year_mod),
          olsrr::ols_press(int.all_mod),
          olsrr::ols_press(step_c),
          olsrr::ols_press(step_d))
```

```
data.frame(df_e, df_mod, r2, adj_r2, se_e, aic, press)
```

	df_e	df_mod	r2	adj_r2	se_e	aic	press
## 1	1895	1	0.0000000	0.0000000	0.7400362	4242.010	1038.8991
## 2	1880	16	0.7262950	0.7241111	0.3887049	1815.354	289.2403
## 3	1854	42	0.7614822	0.7562076	0.3653954	1606.451	259.4941
## 4	1777	119	0.8283172	0.8169167	0.3166484	1137.056	200.2715
## 5	1866	30	0.7611447	0.7574326	0.3644762	1585.132	256.5791
## 6	1808	88	0.8273494	0.8190415	0.3148056	1085.713	195.3835

## Question: 7

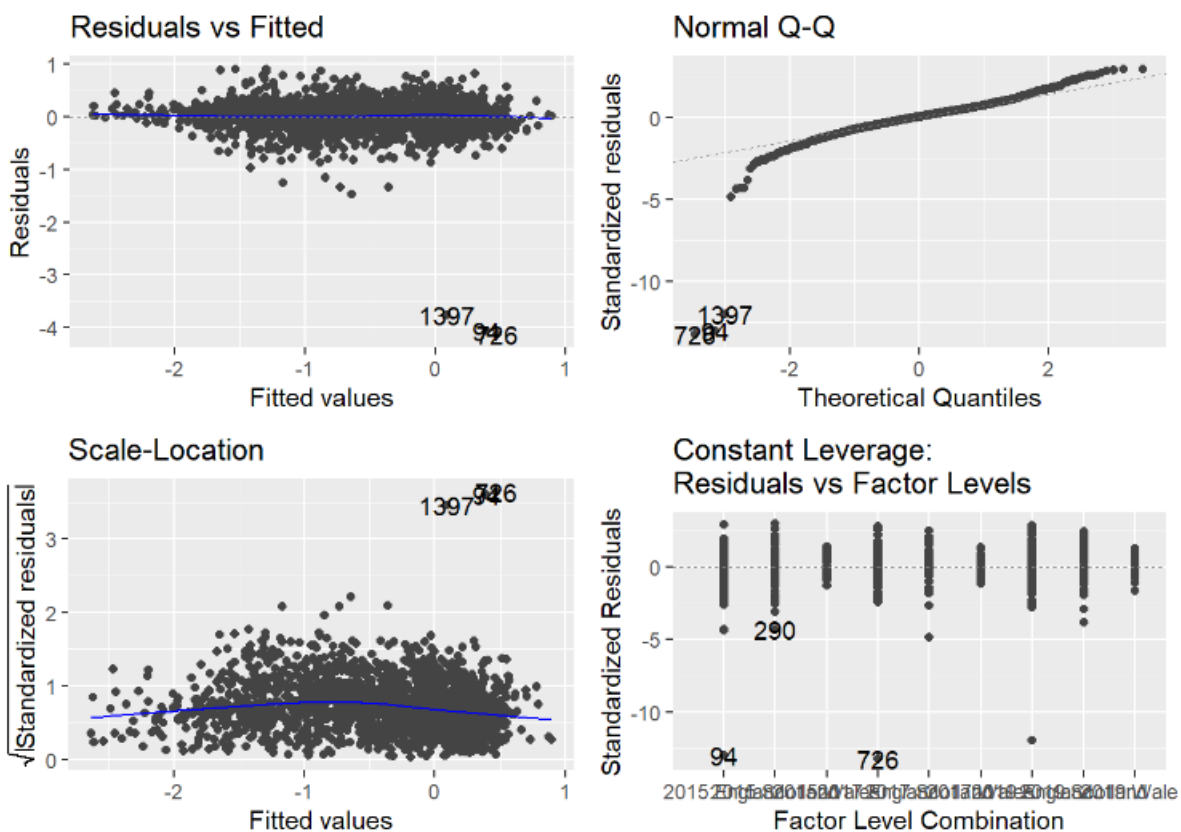
### Solution:

From the above table we observed that the *step\_d* might be the best model depending on the model metrics. However, the full model also perform well to explain the variation of Conservative vote share. The model with interaction can improve a bit to explain the variation of Conservative vote share.

## Question: 8

### Solution:

```
autoplot(step_d)
```



From the above diagnostic plot, we observed that the residual are approximately normally distributed but three observations are clearly outlier.

## Question: 9

### Solution:

```
summary(step_d)$coefficients[c("country_nameScotland", "foreignborn", "country_nameScotland:health.bad"), ]
```

##	Estimate	Std. Error	t value
Pr(> t )			

```
## country_nameScotland      -56.168179  7.4680828 -7.521097 8.507891e-
14
## foreignborn                -1.696258  0.2450373 -6.922450 6.143609e-
12
## country_nameScotland:health.bad  1.228138  0.1361882  9.017947 4.760267e-
19
```

The variables `country_name`, `foreignborn` , and the interaction `country_name` with `health.bad` are highly significant to reduce the variance in the “best” model. The proportion of vote in Scotland is very much less than the England. For *foreignborn*, we observed same scenarios like Scotland. Nevertheless, the interaction between Scotland and `health.bad` have positive association with proportion of vote, i.e, the residence in Scotland with poor health condition are likely to give their vote Conservative party.